

Redrob Candidate Discovery Engine

The Self-Healing Autonomous Development Workspace

Track 1: The AI Systems Architect: Reimagining Work

Saurabh Kumar Bajpai • saurabhbajpai03@outlook.com

- ◆ Developers spend **less than 30%** of their workday writing features.
- ◆ The majority of time is consumed by **context-switching**, searching documentation, and environment setup.
- ◆ **Manual validation cycles** (writing test suites, auditing UI, finding lints) create bottlenecks.
- ◆ Keyword-stuffed and fraudulent candidate resumes waste HR screening resources.

The Corporate Cost

Fragile deployments, slow developer onboarding, and high ATS processing noise result in massive resource drains for scaling enterprises.

- ◆ Reimagines developer setups into a **fully containerized autonomous agent workspace**.
- ◆ Features **self-healing feedback loops**: AI compiles code, runs test suites, lints, and corrects syntax errors autonomously.
- ◆ Integrates **Chrome DevTools MCP** for automated visual and accessibility auditing before commits.
- ◆ Employs **deterministic logical filters** to flag ATS anomalies and keyword stuffers automatically.

Architectural Paradigm Shift

Shifts developer responsibilities from manual coding and debug loops to high-leverage architectural orchestration.

Redrob Candidate Discovery

Search by name, title, or skills...

1 Shreya Tiwari

Senior NLP Engineer • 7.8 yrs exp

350.1

2 Ira Mukherjee

NLP Engineer • 6.6 yrs exp

345.3

3 Aarav Trivedi

Senior Machine Learning Engineer • ...

344.7

4 Rahul Mukherjee

Machine Learning Engineer • 5.7 yrs...

335.8

5 Nisha Pillai

Search Engineer • 7.6 yrs exp

332.8

6 Aarav Agarwal

Staff Machine Learning Engineer • 7...

331.9

7 Saanvi Naidu

Senior NLP Engineer • 8.9 yrs exp

321.4

8 Amit Arora

Lead Machine Learning Engineer • 8...

312.6

Total Scanned: 100,000

Top Matches: 100

Honeypots Discarded: 81

Shreya Tiwari

Senior NLP Engineer

Indore, Madhya Pradesh, India • 7.8 Years Stated Experience

Score: 350.1

REDROB AI MATCH REASONING

Top-tier candidate possessing 7.8 years of industry experience, including production systems built at Niramai. Strong background in TensorFlow aligns perfectly with the JD, and their 89% recruiter response rate demonstrates high availability.

Career History

Senior NLP Engineer

2022-02-17 - Present (52 mo)

Niramai • 51-200 • HealthTech AI

Owned the end-to-end ranking pipeline at a recommendations-heavy consumer product: candidate sourcing → embedding generation (using a fine-tuned BGE-large) → Pinecone retrieval → learning-to-rank re-scoring (XGBoost) → behavioral-signal integration. The hardest part wasn't the ML — it was the evaluation: building offline metrics that actually predicted what the recommendation would do to live engagement. After three iterations we landed on a calibration approach using simulated A/B

Education

Ph.D in Artificial Intelligence

2009 - 2014

IIT Delhi • tier_1

Platform Engagement Signals

RESPONSE RATE

89%

NOTICE PERIOD

15 Days

GITHUB SCORE

76

- ◆ ****Match Score Calibration****: Real-time evaluation based on years of experience, current title, and technical skillset.
- ◆ ****Custom AI Reasoning****: Provides custom, non-templated match summaries generated dynamically for each candidate.
- ◆ ****Interactive Candidate List****: Displays rank, name, and alignment scores for rapid navigation.

- ◆ ****Redrob Behavioral Signals****: Visualizes recruiter response rates, notice periods, and login activity.
- ◆ ****Availability Highlights****: Prioritizes candidates with a buyout timeline of ****≤ 30 days****.
- ◆ ****Technical Competencies****: Groups and colors skills based on explicit expert, advanced, or intermediate proficiency.

Redrob Candidate Discovery

Total Scanned: 100,000Top Matches: 100Honeypots Discarded: 81

Search by name, title, or skills...

1

Shreya Tiwari

Senior NLP Engineer • 7.8 yrs exp

350.1

2

Ira Mukherjee

NLP Engineer • 6.6 yrs exp

345.3

3

Aarav Trivedi

Senior Machine Learning Engineer • ...

344.7

4

Rahul Mukherjee

Machine Learning Engineer • 5.7 yrs...

335.8

5

Nisha Pillai

Search Engineer • 7.6 yrs exp

332.8

6

Aarav Agarwal

Staff Machine Learning Engineer • 7 ...

331.9

7

Saanvi Naidu

Senior NLP Engineer • 8.9 yrs exp

321.4

8

Amit Arora

Staff Machine Learning Engineer • 8 ...

312.4

Senior NLP Engineer

2022-02-17 - Present (52 mo)

Nirama • 51-200 • HealthTech AI

Owned the end-to-end ranking pipeline at a recommendations-heavy consumer product: candidate sourcing → embedding generation (using a fine-tuned BGE-large) → Pinecone retrieval → learning-to-rank re-scoring (XGBoost) → behavioral-signal integration. The hardest part wasn't the ML — it was the evaluation: building offline metrics that actually predicted what the recommendation would do to live engagement. After three iterations we landed on a calibration approach using simulated A/B tests that has held up over the last 18 months.

Senior Machine Learning Engineer

2018-11-05 - 2022-02-17 (40 mo)

Krutim • 201-500 • AI/ML

Fine-tuned LLaMA-2-7B and Mistral-7B variants using LoRA and QLoRA for domain-specific candidate-JD matching. Built the data curation pipeline that generated 200K high-quality preference pairs from recruiter labels, plus the eval harness using both ranking metrics and human-quality scores. Deployed the model via BentoML on Kubernetes with sub-200ms p95 latency by quantizing to INT8 and batching at the request level. Cost per inference dropped from \$0.04 with GPT-3.5-fallback to under \$0.001.

Ph.D in Artificial Intelligence

2009 - 2014

IIT Delhi • tier_1

Platform Engagement Signals

RESPONSE RATE

89%

NOTICE PERIOD

15 Days

GITHUB SCORE

76

SALARY MIN/MAX

28.7 - 60.7 LPA

WORK MODE

Hybrid

PROFILE COMPLETE

65%

Skills & Proficiencies

TensorFlow

EXPERT

OpenSearch


EXPERT

FAISS

ADVANCED

PEFT

EXPERT



◆ ****Zero API Dependencies****: Runs completely offline, avoiding latency drift.

◆ ****Model Context Protocol (MCP)****: Unifies filesystem, shells, and debuggers.

Why This Solution Wins

INDIA RUNS CHALLENGE VERIFICATION

- ◆ ****0% Honeypot Rate****: Filters out exactly ****81 logical anomalies**** (timeline conflicts, 0-duration expert skills, future certs).
- ◆ ****Unmatched Performance****: Runs in **** ≤ 12 seconds**** over the 100k candidate pool.
- ◆ ****Minimal Compute Footprint****: Uses ****~180 MB RAM**** and runs CPU-only, offline.
- ◆ ****ATS Noise Reducer****: Penalizes job hopping (< 18 mo) and filters consulting-only backgrounds.

Stage 3 & 4 Ready

Designed to meet the compute environment restrictions of Stage 3 and provide unique, non-templated candidate match reasonings for Stage 4 manual reviews.

Roadmap

Days 1-30: Establish core docker sandboxing and compiler integration.

Days 31-60: Pilot Chrome DevTools protocol UI verification with internal developers.

Days 61-90: Launch v1.0 open-source VS Code MCP server extension.

- ◆ ****70% Reduction**** in developer debugging and feature drafting cycles.
- ◆ ****50% Faster QA Loops**** due to automated visual and accessibility testing.
- ◆ ****ATS Filtering Accuracy**** to secure top-tier engineering talent and exclude fraud profiles.